

オペレーティングシステム特論

Advanced Operating Systems

#1 (2008/10/3)

Fall 2008, Mon. 9:00-10:30, Rm. W831

Instructor: Takuo Watanabe (渡部卓雄)

e-mail: aos08@psg.cs.titech.ac.jp

ext. 3690, office W8E-805

<http://www.psg.cs.titech.ac.jp/aos>

1

Outline

- ➔ Introduction
 - About This Course
 - Course Administtrivia
- High-Level Language VMs (1)
 - Implementing Simple Stack VM

2

About This Course

- Goal: to give students a better understanding of Virtual Execution Environments (仮想実行環境) and their applications.

3

Topics to be Covered (1)

- Process Virtual Machines
 - High-Level Language VMs (HLL-VMs)
 - JVM, Parrot, YARV, etc.
 - WAM, CAM, SECD
 - Other Process-Level VMs
 - Jail, Zap, etc.
- System Virtual Machines
 - VMM, Paravirtualization, Microkernel
 - Xen, Denali, Exokernel, etc.

4

Topics to be Covered (2)

- Applications of VMs
 - ex. Secure Computing
 - Terra, Revirt, VM-based Introspection for Intrusion Detection, Collapsar
 - Honeypot, HoneyNet Project
 - Monitored Execution

5

Prerequisites

- Basic knowledge on:
 - Programming Languages (esp. C)
 - Operating Systems
 - Computer Architectures
 - Logic (First Order Logic)

6

Administrativa

- Workload:
 - lectures
 - present at least one paper on VEE
 - report (TBA) on VEE
- Grading:
 - 50% presentation
 - 50% report

7

Administrativa

- 60 min. lecture
- 30 min. presentation
 - 20 min. talk + 10 min. discussion
 - Presentation may be either in English or Japanese.
 - Presentation materials (slides and handouts) should be written in English.
 - Reading list is on the course web site. Pick up 3 papers and send me their titles with your preference order until Oct. 9.

8

Presentation Guidelines

- Make sure that you understand the paper.
- Prepare some slides.
 - 1 slide \approx 2 minutes
- Rehearse at least once.
- Split the presentation in three parts.
 - Overview of the paper and explanation of the problem
 - Discussion of key contribution
 - Critique

9

Calendar (1)

- 10/3 lecture (1) (introduction)
- 10/10 lecture (2)
- 10/17 lecture (3) + presentation (1)
- 10/31 lecture (4) + presentation (2)
- 11/7 lecture (5) + presentation (3)
- 11/14 lecture (6) + presentation (4)
- 11/21 lecture (7) + presentation (5)
 - (11/28 no class)
- 12/5 lecture (8) + presentation (6)

10

Calendar (2)

- 12/12 lecture (9) + presentation (7)
- 12/19 lecture (10) + presentation (8)
- 1/9 lecture (11) + presentation (9)
- 1/23 lecture (12) + presentation (10)
 - or presentations (10-12)
- 1/30 presentations (11-13)
 - or presentations (13-15)
- (2/6 presentations)

11

Lecture Plan (1)

- Introduction
- High Level Language VMs (1)
 - Interpreters
 - Simple Bytecode Interpreter
 - Various Improvements of Interpreters
 - Threaded Code, Stack Caching, Superoperators

12

Lecture Plan (2)

- High-Level Language VMs (2)
 - Java VM
 - Bytecode Verification, JIT
 - VM for other languages: CAM, WAM, etc.
- Runtime Systems
 - Memory Management
- Process VMs
- System VMs
 - VMM, Paravirtualization

13

Books

- Virtual Machines
 - Iain D. Craig
 - Springer, 2006 (1-85233-969-1).
- Virtual Machines: Versatile Platforms for Systems and Processes
 - James E. Smith & Ravi Nair
 - Morgan Kaufmann, 2005 (1-55860-910-5).

14

Outline

- Introduction
 - About This Course
 - Course Administria
- ➔ High-Level Language VMs (1)
 - Implementing Simple Stack VM in C


15

Virtual Machine

- Classic Definition
 - "an efficient, isolated duplication of a real machine"
 - G. J. Popek & R. P. Goldberg, "Formal Requirements for Virtualizable Third Generation Architectures", Comm. ACM, 17(7), 1974.
- Current
 - a software that provides virtual software execution environment
 - not restricted to real machines
 - Taxonomy
 - System VM
 - Process VM

16

Machine Interfaces



Application Binary Interface (ABI)


= System Calls + User ISA

Application Program Interface (API)

= System Calls + Library Calls + User ISA


17

System Virtual Machine



18

Process Virtual Machine



19

High Level Language VM

- An interpreter for virtual (abstract) machine instructions
 - a kind of process VMs
- Constructions
 - Execution Mechanisms
 - Stack Machines
 - JVM, .Net CLR, etc.
 - Register Machines
 - CAM, Parrot, etc.
 - Instruction Representation
 - Bytecode
 - Wordcode, VLIW-like code, etc.

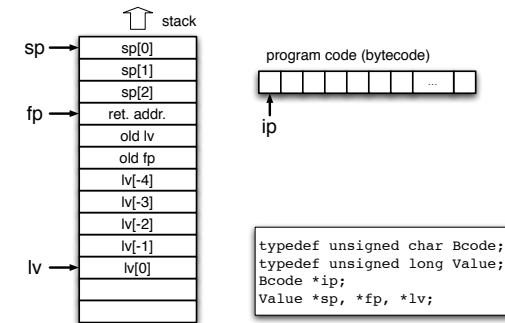
20

Example: A Simple VM in C

- Bytecode Machine
 - Java-like instruction set
- Stack Machine
 - Function arguments and local variables are on the value stack.
 - 32-bits value
 - integer
 - function pointer
- Source code can be obtained from the course website.

21

VM Structure



22

Instructions

- nop, dup
- ifeq, ifne, iflt, ifle, ifgt, ifge
- giload, gaload, gistore, gastore
- iload, aload, istore, astore
- bipush, sipush, iconst_m1, iconst_0, iconst_1, iconst_2
- iadd, isub, imul, idiv, irem
- icall, acall, ticall, tacall, sicall, sacall
- iret, aret
- goto, halt

23

The Main VM Function

```
void vm (Bcode code[], size_t codelen,
        Value *vs_base, size_t vs_size,
        value gv[], size_t gv_size,
        vm_mode mode) {
    Bcode *ip = code;      // instruction pointer
    Value *sp = vs_base;   // stack pointer
    Value *fp = sp;        // frame pointer
    Value *lv = sp;        // local vars base pointer

    for (;;) {
        switch (*ip) {
            // giant switch implementing instructions
        }
    }
}
```

24

Fetch-Decode-Execute (Giant Switch)

```
for (;;) {
    switch (*ip) {
        case i_nop:
            ip++;
            break;
            :
        case i_iadd:
            sp[1] = sp[1] + sp[0];
            sp++;
            ip++;
            break;
            :
    }
}
```

25

Implementing Instructions

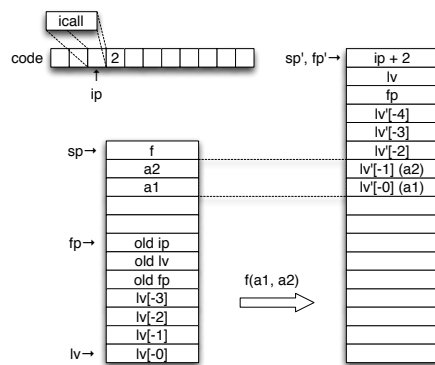
```
iadd:
    sp[1] = sp[1] + sp[0];
    sp++;
    ip++;
```

```
iload:
    sp--;
    sp[0] = lv[-(int)ip[1]]
    ip += 3;
```

```
goto:
    ip += ((signed char)ip[1] << 8 | ip[2]) + 3;
```

26

Function Call




Function Call Implementation

```
icall:
    Function f = (Function)sp[0];
    int nargs = (int)ip[1];
    Value *newfp = sp - f->nlv - 2;
    newfp[0] = (Value)(ip + 2);
    newfp[1] = (Value)lv;
    newfp[2] = (Value)fp;
    lv = sp + nargs;
    sp = fp = newfp;
    ip = f->code;
```

28

Tail-Call



29


Tail-Call Implementation

```

tcall:
    Function f = (Function)sp[0];
    int nargs = (int)ip[1];
    Value ip_save = fp[0];
    Value lv_save = fp[1];
    Value fp_save = fp[2];
    Value *arg = sp + nargs;
    int i;
    for (i = 0; i < nargs ; i++)
        lv[-i] = arg[-i];
    fp = lv - nargs - f->nlv - 3;
    sp = fp;
    fp[0] = ip_save;
    fp[1] = lv_save;
    fp[2] = fp_save;
    ip = f->code;
    
```

30

Threaded Code



31

Direct Threaded Code (DTC)

```

iadd:
    sp[1] = sp[1] + sp[0];
    sp++;
    ip++;
    goto tcode[ip];
    
```

32

Labels as Value Extension in GCC

```
void *laddr;  
  
laddr = &&label;  
  
goto *laddr;  
  
label:  
....
```

33

DTC in GCC

```
#define NEXT goto **ip  
  
typedef void *Tcode;  
  
Tcode tcode[] = ...; // Program in DTC  
Tcode *ip = tcode;  
NEXT;  
  
l_nop:  
    ip++;  
    NEXT;  
    :  
l_iadd:  
    sp[1] = sp[1] + sp[0];  
    sp++;  
    ip++;  
    NEXT;  
    :
```

34

vmbench.c

- Result
 - svm (Giant Switch)
 - 14.96 sec (PPC), 4.27 sec (IA32)
 - tvn (DTC)
 - 10.85 sec (PPC), 4.78 sec (IA32)
- Environment (OS X 10.4.10, GCC 4.0.1):
 - PPC: 1.25GHz PPC G4 , 2GB RAM
 - IA32: 2.33GHz Core2 Duo, 2GB RAM

35

Summary

- Course Guidance
- Implementing Simple Stack VM in C

36